

Predicting Hypertension

In-depth Analysis with Machine Learning

Choosing Algorithms

Algorithms were chosen based on the properties of the data and target variable. The data from the National Health Interview Survey and is labeled and contains the target variable, hypertension, making this a supervised learning problem. The target variable has two categories, yes/no, making this a binary classification problem. Five algorithms were chosen for this problem type, including K-Nearest Neighbors (KNN), Logistic Regression, Random Forest Classifier, Adaptive Boosting (AdaBoost) and Extreme Gradient Boosting (XGBoost).

Choosing Evaluation Metrics

Evaluation metrics were chosen based on the properties of the data and the implication of false predictions. In this case, the data is imbalanced, with 39% of responses positive for hypertension. For this reason, F1 score was chosen over accuracy.

The F1 score was also appropriate due to a difference in the cost of false positives (predicting hypertension when the condition does not exist) and false negatives (predicting no hypertension when the condition does exist), making it better to focus on the true positive rate, also known as sensitivity or recall, and precision, or how "precise" the classifier is when predicting positive instances. Since the F1 score is the weighted

average of precision and recall, if the F1 score is high, both precision and recall of the classifier indicate good results.

Initial Results

To establish a baseline for the results, the first four algorithms were run with default parameters on the dataset produced after data wrangling. Parameters were changed only to resolve issues, such as increasing the number of iterations to resolve a convergence warning for logistic regression. The results were:

	Positive F1	Negative F1
KNN	0.38	0.82
Logistic Regression	0.45	0.84
Random Forest	0.49	0.83
AdaBoost	0.47	0.88

Improving Results – Additional Data Wrangling and Feature Selection

The first approach at improving these results was to review the data again. Based on re-examining usable data, one feature was dropped and missing data for two features was removed, resulting in a dataset of 484,070 rows and 39 columns. Feature importance was reviewed but values were inconsistent among models. The models were re-run on this set, showing improvements on the positive F1 score:

	Positive F1	Negative F1
KNN	0.43 +11%	0.74
Logistic Regression	0.53 +17%	0.78
Random Forest	0.54 +10%	0.84

AdaBoost	0.56 +19%	0.78
----------	-----------	------

The second approach at improving these results was to test the impact of scaling the data. The scaling preprocessing method was tested with KNN, improving the positive F1 score to 0.47. Since this result was not as good as other models, scaling was not implemented.

Improving Results – Tuning Algorithms, Additional Training/Testing

The third approach at improving these results was to adjust the settings on the two top-performing algorithms and run them through more training and testing iterations. This was achieved with hyperparameter tuning and cross validation via grid search and random search methods. Multiple rounds of different values for parameters were performed, until the results leveled off here:

	Positive F1	Negative F1
Logistic Regression	0.54 +1%	0.78
AdaBoost	0.58 +3%	0.79

Final Results - XGBoost

Having adjusted the data and tuned the models, the final approach was to try a different model. XGBoost was chosen for its consistent outperformance of other models. Using multiple rounds of hyperparameter tuning and cross validation, XGBoost produced the highest positive next-highest negative F1 scores:

	Positive F1	Negative F1
XGBoost	0.60	0.78

XGBoost also produced a model with the highest precision, recall and specificity:

- If the model predicts hypertension, this result will be correct 68% of the time (precision)
- If the subject has hypertension, the model predicts this accurately 54% of the time (sensitivity, recall or true positive rate)
- If the subject does not have hypertension, the model predicts this accurately 83% of the time (specificity or true negative rate)

Summary

The results are less than what was hoped for, and while this model is specific and somewhat precise, it is not sensitive. Reasons for these results could be due biases in the data. Not all candidates for the survey respond, creating selection bias. When the survey is conducted, not all questions are asked, and not questions are answered, creating a bias among the data. In addition, the information is self-reported rather than obtained from verified sources, creating opportunity for inaccuracies in the answers provided.

Fortunately, despite the low sensitivity, the model can still provide benefit. Tests with high specificity provide value when the result is positive, as they can be used for ruling in patients who have a certain disease. With the current results, the usefulness of the model is based less on its ability to predict hypertension, but more on its function as a screening tool for identifying people at risk for hypertension. This model could be used in community outreach efforts to raise awareness about hypertension and identify people who should follow up with their primary care provider to have their blood pressure tested. It could also be used to flag patients in the healthcare setting for a focused assessment of hypertension.

Future Improvements

These results could be potentially be improved with additional efforts on feature selection, such as using methods designed for this purpose available in Scikit-Learn (the SelectKBest method) and pruning features based on their importance in tree-based algorithms.

